

STA663 Final Project

Infinite Latent Feature Models and the Indian Buffet Process

Dipesh Gautam

April 29, 2015

1 Indian Buffet Process (IBP)

The Indian Buffet is an adaptation of Chinese Buffet Process where each object instead of being associated with a single latent class can be associated with multiple classes. This is particularly useful when each object has multiple latent features and by associating objects with a single class we cannot partition them into homogeneous subsets.

In the Indian buffet process, N customers enter a restaurant one after another. Each customer encounters a buffet consisting of infinitely many dishes arranged in a line. The first customer starts at the left of the buffet and takes a serving from each dish, stopping after a $\text{Poisson}(\alpha)$ number of dishes. The i th customer moves along the buffet, sampling dishes in proportion to their popularity, taking dish k with probability $\frac{m_k}{i}$, where m_k is the number of previous customers who have sampled that dish. Having reached the end of all previous sampled dishes, the i th customer then tries a $\text{Poisson}(\frac{\alpha}{i})$ number of new dishes. Which customer chose which dishes is indicated using a binary matrix \mathbf{Z} with N rows and infinitely many columns (corresponding to the infinitely many selection of dishes), where $z_{ik} = 1$ if the i th customer sampled k th dish.

IBP can be used as a prior in models for unsupervised learning. An example of which is presented in the paper by Griffiths and Ghahramani, where IBP is used as a prior in linear-Gaussian binary latent feature model.

2 Algorithm

- Gamma prior for α

$$\alpha \sim \text{Gamma}(1, 1)$$

- Prior on \mathbf{Z} is obtained by IBP as:

$$P(z_{ik} = 1 | \mathbf{z}_{-i,k}) = \frac{n_{-i,k}}{N}$$

- Likelihood is given by

$$P(X | Z, \sigma_X, \sigma_A) = \frac{1}{(2\pi)^{ND/2} (\sigma_X)^{(N-K)D} (\sigma_A)^{KD} (|Z^T Z + \frac{\sigma_X^2}{\sigma_A^2} I|)^{D/2}} \exp\left\{-\frac{1}{2\sigma_X^2} \text{tr}(X^T (I - Z(Z^T Z + \frac{\sigma_X^2}{\sigma_A^2} I)^{-1} Z^T) X)\right\} \quad (1)$$

After we have the likelihood and the prior given by IBP,

- full conditional posterior for \mathbf{Z} can be calculated as:

$$P(z_{ik} | X, Z_{-(i,k)}, \sigma_X, \sigma_A) \propto P(X | Z, \sigma_X, \sigma_A) * P(z_{ik} = 1 | \mathbf{z}_{-i,k})$$

To sample the number of new features for observation i , we use a truncated distribution, computing probabilities for a range of values $K_1^{(i)}$ up to an upper bound (say 4). The prior on number of features is given by $\text{Poisson}(\frac{\alpha}{N})$. Using this prior and the likelihood, we sample the number of new features.

- Full conditional posterior for α is given by:

$$P(\alpha|Z) \sim \text{Gamma}(1 + K_+, 1 + \sum_{i=1}^N H_i)$$

- For σ_X and σ_A , we use MH algorithm as follows:

$$\epsilon \sim \text{Uniform}(-.05, .05) \quad (2)$$

$$\sigma_X^* = \sigma_X + \epsilon \quad (3)$$

$$(4)$$

Accept this new σ_X with probability given by:

$$AR = \min\{1, \frac{\text{Likelihood}(X|\sigma_X, \dots)}{\text{Likelihood}(X|\sigma_X, \dots)}\}$$

Where AR is the acceptance ratio. We use similar algorithm to sample σ_A

3 Profiling

We profiled the code using *cProfile* to figure out the bottleneck. The result is shown in *profiler.txt*. We see that most of the computational time is spent on calculating the *log likelihood(l)* and matrix inversion. Due to this fact, one of the first things we looked at were ways to reduce computation time for likelihood and/or inverse calculation.

profiler.txt

2075808 function calls in 10.873 seconds

Ordered by: internal time

ncalls	totttime	percall	cumtime	percall	filename:lineno:function
154080	3.985	0.000	3.985	0.000	{method 'dot' of 'numpy.ndarray' objects}
30816	1.540	0.000	9.411	0.000	<ipython-input-144-816e3f6a3e53>:211
30816	0.770	0.000	1.353	0.000	linalg.py:455:inv
1	0.726	0.726	10.873	10.873	<ipython-input-145-4efe9a6e9287>:1:sampler
30816	0.542	0.000	1.071	0.000	linalg.py:1679:det
61632	0.382	0.000	0.963	0.000	numeric.py:2125:identity
61632	0.367	0.000	0.580	0.000	twodim_base.py:190:eye
30816	0.332	0.000	0.332	0.000	{method 'trace' of 'numpy.ndarray' objects}
86971	0.267	0.000	0.267	0.000	{numpy.core.multiarray.zeros}
61632	0.183	0.000	0.314	0.000	linalg.py:139:_commonType
122449	0.171	0.000	0.171	0.000	{numpy.core.multiarray.array}
20964	0.135	0.000	0.135	0.000	{method 'reduce' of 'numpy.ufunc' objects}
30816	0.114	0.000	0.114	0.000	{method 'astype' of 'numpy.ndarray' objects}
92448	0.113	0.000	0.251	0.000	numeric.py:394:asarray
30816	0.107	0.000	0.107	0.000	{method 'astype' of 'numpy.generic' objects}
71965	0.104	0.000	0.104	0.000	{max}
15000	0.092	0.000	0.092	0.000	{numpy.core.multiarray.concatenate}
61632	0.089	0.000	0.147	0.000	linalg.py:209:_assertNdSquareness
30000	0.084	0.000	0.166	0.000	shape_base.py:8:atleast_1d
61632	0.073	0.000	0.086	0.000	linalg.py:198:_assertRankAtLeast2
184896	0.071	0.000	0.071	0.000	{issubclass}
30816	0.067	0.000	0.500	0.000	fromnumeric.py:1233:trace
20964	0.053	0.000	0.231	0.000	fromnumeric.py:1631:sum
123264	0.052	0.000	0.080	0.000	linalg.py:1111:isComplexType
30816	0.052	0.000	0.052	0.000	linalg.py:101:get_linalg_error_extobj
15000	0.049	0.000	0.307	0.000	shape_base.py:230:hstack
30816	0.048	0.000	0.129	0.000	linalg.py:106:_makearray

10980	0.040	0.000	0.040	0.000 {method 'uniform' of 'mtrand.RandomState' objects}
30816	0.037	0.000	0.037	0.000 linalg.py:219_assertNoEmpty2d
61632	0.036	0.000	0.048	0.000 linalg.py:124_realType
30000	0.030	0.000	0.063	0.000 numeric.py:464asanyarray
20964	0.026	0.000	0.026	0.000 {isinstance}
121632	0.024	0.000	0.024	0.000 {len}
5000	0.019	0.000	0.019	0.000 {sum}
20964	0.017	0.000	0.152	0.000 _methods.py:31_sum
61732	0.015	0.000	0.015	0.000 {min}
61632	0.012	0.000	0.012	0.000 {method 'get' of 'dict' objects}
20000	0.011	0.000	0.011	0.000 {math.factorial}
15152	0.011	0.000	0.011	0.000 {range}
30816	0.010	0.000	0.010	0.000 {getattr}
30000	0.009	0.000	0.009	0.000 {method 'append' of 'list' objects}
30816	0.007	0.000	0.007	0.000 {method '__array_prepare__' of 'numpy.ndarray' objects}
1	0.002	0.002	0.005	0.005 <ipython-input-143-ed70069a6371>:2sampleIBP
100	0.000	0.000	0.000	0.000 {method 'poisson' of 'mtrand.RandomState' objects}
50	0.000	0.000	0.000	0.000 {method 'gamma' of 'mtrand.RandomState' objects}
4	0.000	0.000	0.000	0.000 {numpy.core.multiarray.copyto}
4	0.000	0.000	0.000	0.000 {numpy.core.multiarray.empty}
4	0.000	0.000	0.000	0.000 numeric.py:141ones
1	0.000	0.000	0.000	0.000 {method 'seed' of 'mtrand.RandomState' objects}
1	0.000	0.000	10.873	10.873 <string>:1<module>
1	0.000	0.000	0.000	0.000 {method 'disable' of '_lsprof.Profiler' objects}

3.1 Improved matrix Inversion

We tried the matrix inversion method described in Griffiths and Ghahramani(2005, eq 51-54), where the method reduces the runtime by allowing us to perform rank one updates instead when only one value is changed. We implemented the algorithm and were able to speed up the process as shown in Table 1.

Table 1: Comparison of matrix inverse methods

	Time
linalg.inverse	0.000078
calcInverse	0.000036

Even though we were able to improve the performance, due to some numerical errors, we were not able to obtain a stationary MCMC chain using this method. This could be achieved by spending some more time on it but due to lack of time, we had to abandon this method and move on.

3.2 Improved likelihood function

Table 2: Runtime Comparision

	Total Time
Initial Code	6.495043
Improved ll	6.099027
Cythonized	5.953709

Figure 1: Original Features and First four simulated objects




Figure 2: Features Detected after MCMC and First four recreated objects




Figure 3: Traceplots for σ_X , σ_A and α after burn-in




Figure 4: Distribution of Kplus




Table 3: Runtime Comparision

	F1	F2	F3	F4
1st image	0	1	0	0
2nd image	1	1	0	0
3rd image	1	1	0	1
4th image	1	1	0	0